

PRACTICAL 2 - Implementing Map-Reduce Program for Word Count problem.

SERVICES TO START-

HDFS

HBASE

If name error occurs , restart HDFS

Steps to Perform the Practical:

Step 1: Open Cloudera

- Start your Cloudera environment and log in.

Step 2: Go to Eclipse

- Launch Eclipse from the Cloudera environment.

Step 3: Create a new Java Project

- Go to `File > New > Java Project`.

Step 4: Name the Project and Configure Libraries

- Give the project a name, for example, "WordCount."
- Click "Next" and navigate to the "Libraries" tab.
- Click "Add External JARs."

Step 5: Add Hadoop JARs

- Navigate to the Hadoop library directory, typically located at `/usr/lib/hadoop`.
- Select all the JAR files and click "OK."

Step 6: Add Hadoop Client JARs

- Click "Add External JARs" again.
- Navigate to `/usr/lib/hadoop/client` and select all the JAR files.
- Click "OK."

Step 7: Add Hadoop Client 0.20 JARs

- Once more, click "Add External JARs."
- Navigate to `/usr/lib/hadoop/client0.20` and select all the JAR files.
- Click "OK."

Step 8: Finish Project Configuration

- Click "Finish" to complete the project setup.

Step 9: Create a New Java Class

- Under the newly created Java project, right-click on `src > New > Class`.
- Give the class the same name as the project, e.g., "WordCount," and click "Finish."

Step 10: Write Word Count Code

- Write the Java code for WordCount in the newly created class.

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }
}
```

```

    }

    public static class IntSumReducer
        extends Reducer<Text,IntWritable,Text,IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values,
            Context context
        ) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
            context.write(key, result);
        }
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(TokenizerMapper.class);
        job.setCombinerClass(IntSumReducer.class);
        job.setReducerClass(IntSumReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

- Save the file using `Ctrl + S`.

Step 11: Navigate to the project name, right-click on it, then choose "Export" > "Java" > "JAR file" > "Export to JAR file". Click "Next," set the path for the JAR file (e.g., cloudera), and provide a name matching the project name. Finally, click "Finish" to complete the export process.

OPEN CLOUDERA TERMINAL

Step 1: Check existing files in HDFS

```

-----
[cloudera@quickstart ~]$ hdfs dfs -ls
Found 2 items
drwx----- - cloudera cloudera      0 2024-01-10 03:09 .staging
-rw-r--r--  1 cloudera cloudera     81 2024-01-10 02:59 inputdirectory
-----

```

Step 2: Create a directory mapreduce in HDFS

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -mkdir /mapreduce
```

Step 3: Check the contents of the root directory in HDFS

```
[cloudera@quickstart ~]$ hdfs dfs -ls /
Found 8 items
drwxrwxrwx - hdfs supergroup          0 2017-10-23 10:29 /benchmarks
drwxr-xr-x - hbase supergroup          0 2024-02-16 03:36 /hbase
drwxr-xr-x - hdfs supergroup          0 2024-01-06 06:52 /inputdirectory
drwxr-xr-x - hdfs supergroup          0 2024-02-18 03:08 /mapreduce
drwxr-xr-x - solr solr                 0 2017-10-23 10:32 /solr
drwxrwxrwt - hdfs supergroup          0 2024-01-06 06:56 /tmp
drwxr-xr-x - hdfs supergroup          0 2017-10-23 10:31 /user
drwxr-xr-x - hdfs supergroup          0 2017-10-23 10:31 /var
```

Step 4: Change permissions of the mapreduce directory to 777

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -chmod -R 777 /mapreduce
[cloudera@quickstart ~]$ hdfs dfs -ls / mapreduce
Found 8 items
drwxrwxrwx - hdfs supergroup          0 2017-10-23 10:29 /benchmarks
drwxrwxrwx - hbase supergroup          0 2024-02-16 03:36 /hbase
drwxrwxrwx - hdfs supergroup          0 2024-01-06 06:52 /inputdirectory
drwxrwxrwx - hdfs supergroup          0 2024-02-18 03:08 /mapreduce
drwxrwxrwx - solr solr                 0 2017-10-23 10:32 /solr
drwxrwxrwt - hdfs supergroup          0 2024-01-06 06:56 /tmp
drwxrwxrwx - hdfs supergroup          0 2017-10-23 10:31 /user
drwxrwxrwx - hdfs supergroup          0 2017-10-23 10:31 /var
```

Step 5: Create a local file /home/cloudera/input.txt

```
[cloudera@quickstart ~]$ cat > /home/cloudera/input.txt
myself OM Panchal
om panchal
^Z
[1]+  Stopped                  cat > /home/cloudera/input.txt
```

Step 6: Verify the content of the local file

```
[cloudera@quickstart ~]$ cat > /home/cloudera/input.txt  
myself OM Panchal  
om panchal  
^Z  
[1]+  Stopped                  cat > /home/cloudera/input.txt
```

Step 7: Upload the local file to HDFS under the directory

```
[cloudera@quickstart ~]$ sudo -u hdfs hadoop fs -put /home/cloudera/input.txt /mapreduce
```

Step 8: Verify the uploaded file in HDFS

```
[cloudera@quickstart ~]$ hdfs dfs -ls /mapreduce  
Found 1 items  
-rw-r--r--  1 hdfs supergroup      29 2024-02-18 03:15 /mapreduce/input.txt
```

Step 9: Run the WordCount job

```
[cloudera@quickstart ~]$ hadoop jar /home/cloudera/WordCount.jar WordCount /mapreduce/input.txt /mapreduce/output
```

```

24/02/18 03:19:20 INFO mapreduce.Job: map 100% reduce 0%
24/02/18 03:19:38 INFO mapreduce.Job: map 100% reduce 100%
24/02/18 03:19:39 INFO mapreduce.Job: Job job_1708253660300_0001 completed successfully
24/02/18 03:19:40 INFO mapreduce.Job: Counters: 49

```

File System Counters

```

FILE: Number of bytes read=77
FILE: Number of bytes written=294519
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=145
HDFS: Number of bytes written=39
HDFS: Number of read operations=6
HDFS: Number of large read operations=0
HDFS: Number of write operations=2

```

Job Counters

```

Launched map tasks=1
Launched reduce tasks=1
Data-local map tasks=1
Total time spent by all maps in occupied slots (ms)=12290048
Total time spent by all reduces in occupied slots (ms)=7835136
Total time spent by all map tasks (ms)=24004
Total time spent by all reduce tasks (ms)=15303
Total vcore-milliseconds taken by all map tasks=24004
Total vcore-milliseconds taken by all reduce tasks=15303
Total megabyte-milliseconds taken by all map tasks=12290048
Total megabyte-milliseconds taken by all reduce tasks=7835136

```

Map-Reduce Framework

```

Map input records=2
Map output records=5
Map output bytes=49
Map output materialized bytes=73
Input split bytes=116
Combine input records=5
Combine output records=5
Reduce input groups=5
Reduce shuffle bytes=73
Reduce input records=5
Reduce output records=5
Spilled Records=10
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=514
CPU time spent (ms)=1120
Physical memory (bytes) snapshot=224993280
Virtual memory (bytes) snapshot=1410572288
Total committed heap usage (bytes)=101449728

```

Step 10: Check the output in HDFS

```

[cloudera@quickstart ~]$ hdfs dfs -ls /mapreduce/output
Found 2 items
-rw-r--r-- 1 cloudera supergroup 0 2024-02-18 03:19 /mapreduce/output/_SUCCESS
-rw-r--r-- 1 cloudera supergroup 39 2024-02-18 03:19 /mapreduce/output/part-r-00000

```

Step 11: View the contents of the output file

```

[cloudera@quickstart ~]$ hdfs dfs -cat /mapreduce/output/part-r-00000
OM 1
Panchal 1
myself 1
om 1
panchal 1

```